


# Characterizing and Improving the Use of Demand-Fetched Caches in GPUs

Wenhao Jia, Kelly A. Shaw\*, Margaret Martonosi  
Princeton University   University of Richmond\*

# GPU Caches: Challenges & Opportunities

- GPUs are particularly suitable for data parallel applications
- Newer GPUs have sizable demand-fetched caches
  - Up to 48 KB L1, around 1 MB L2
- NVIDIA GPU L1 caches have considerable configurability
  - Whole program on/off, size, per-instruction on/off




# CPU Caches vs. GPU Caches

- Caches in CPUs
  - Well-studied, optimize hit-rates to reduce DRAM latency
  - 3C's: compulsory, capacity, conflict
- Caches in GPUs
  - New and less well-studied
  - GPU applications have unique characteristics: large number of threads, computational throughput-oriented, complex memory hierarchy, etc.
  - Need a new way to think about and characterize GPU cache utility



# When Are GPU Caches Helpful?

Rodinia benchmark suite, NVIDIA Tesla C2070



- Class I kernels:** Texture/constant loads only, requests don't use L1 caches
- Class II kernels:** Mainly use shared memory, limited benefits from caching
- Class III kernels:** Use DRAM and thus caches frequently, but see large and unpredictable performance variations from caching



# Research Goals

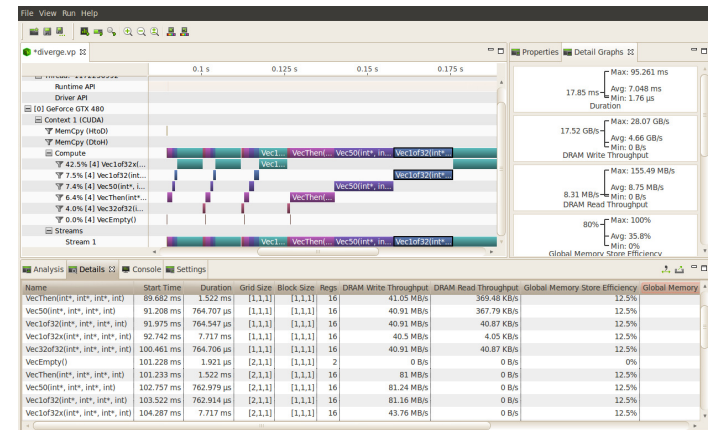
- Characterize performance impact of caches for GPU applications running on real systems
- Present a memory access locality taxonomy to help programmers and compilers analyze cache utility
- Propose and test an automatable, static cache utility analysis algorithm

Prior work has studied data layout and scratchpad optimization, but not caches or their configurability



# Experimental Platform


- Real-system measurements
- NVIDIA Tesla C2070 GPU (Fermi) + Visual Profiler



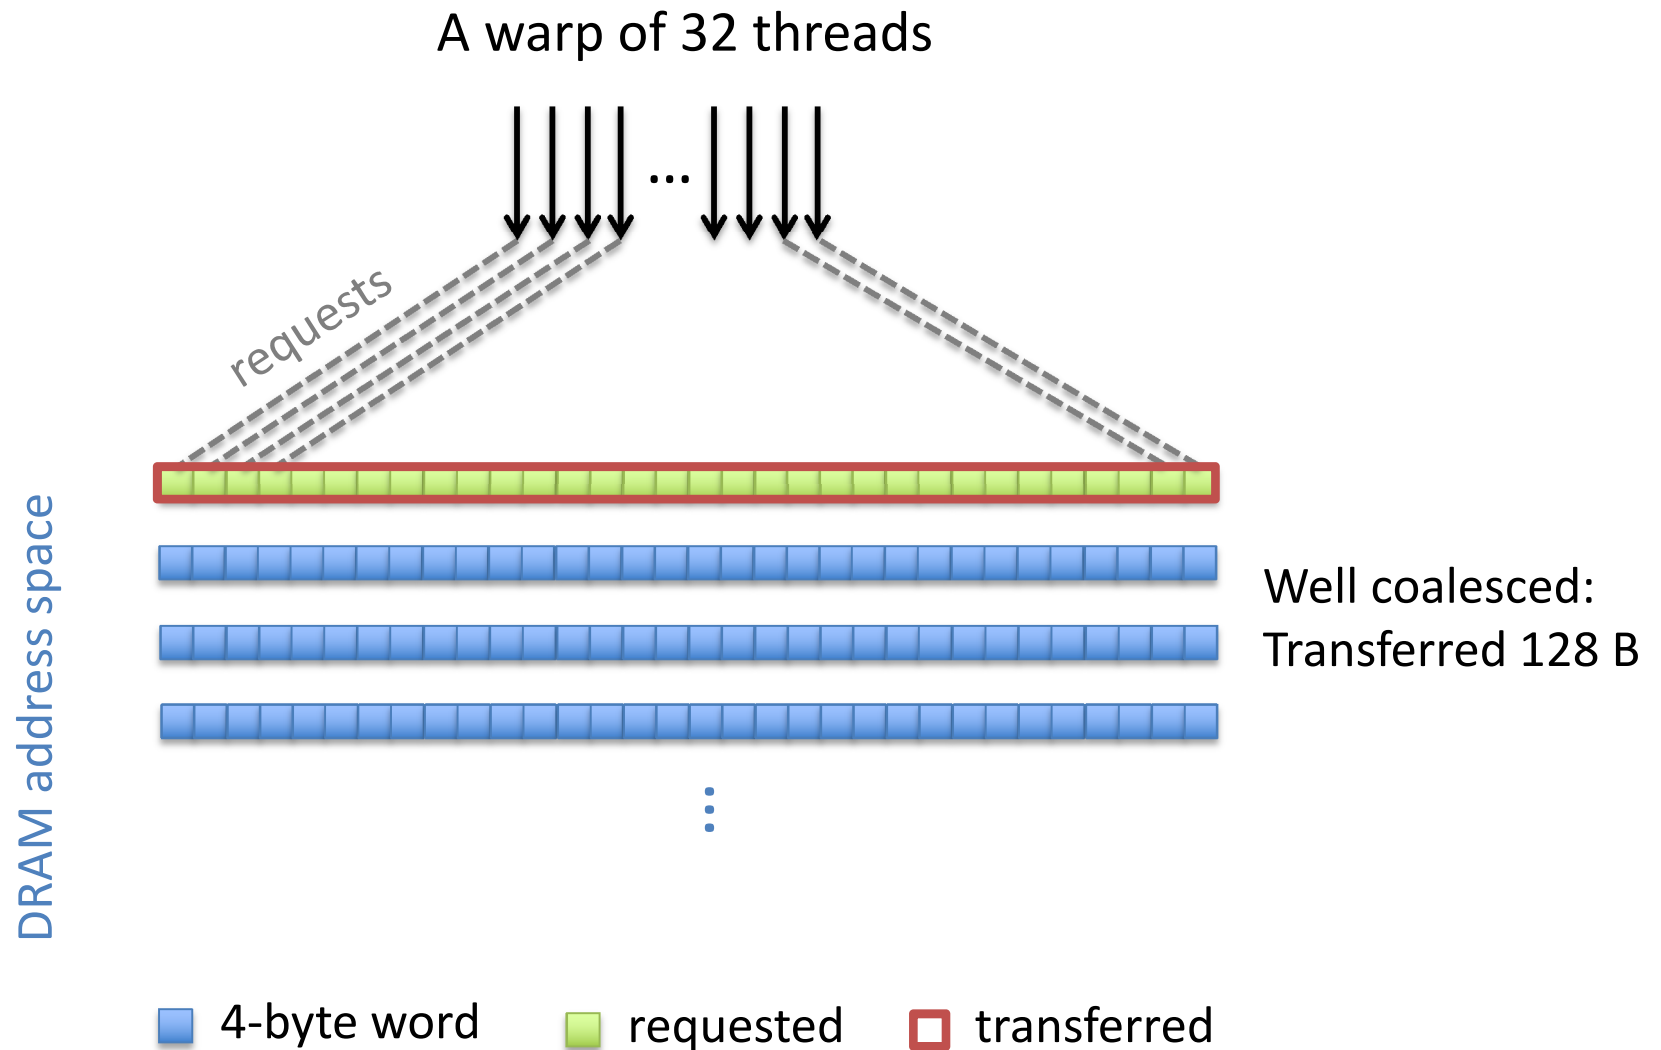
- Rodinia benchmark suite
  - 11 GPGPU applications with 28 kernels

# Background: Warps Hide Memory Latency

Hundreds of concurrent hardware threads in a thread block

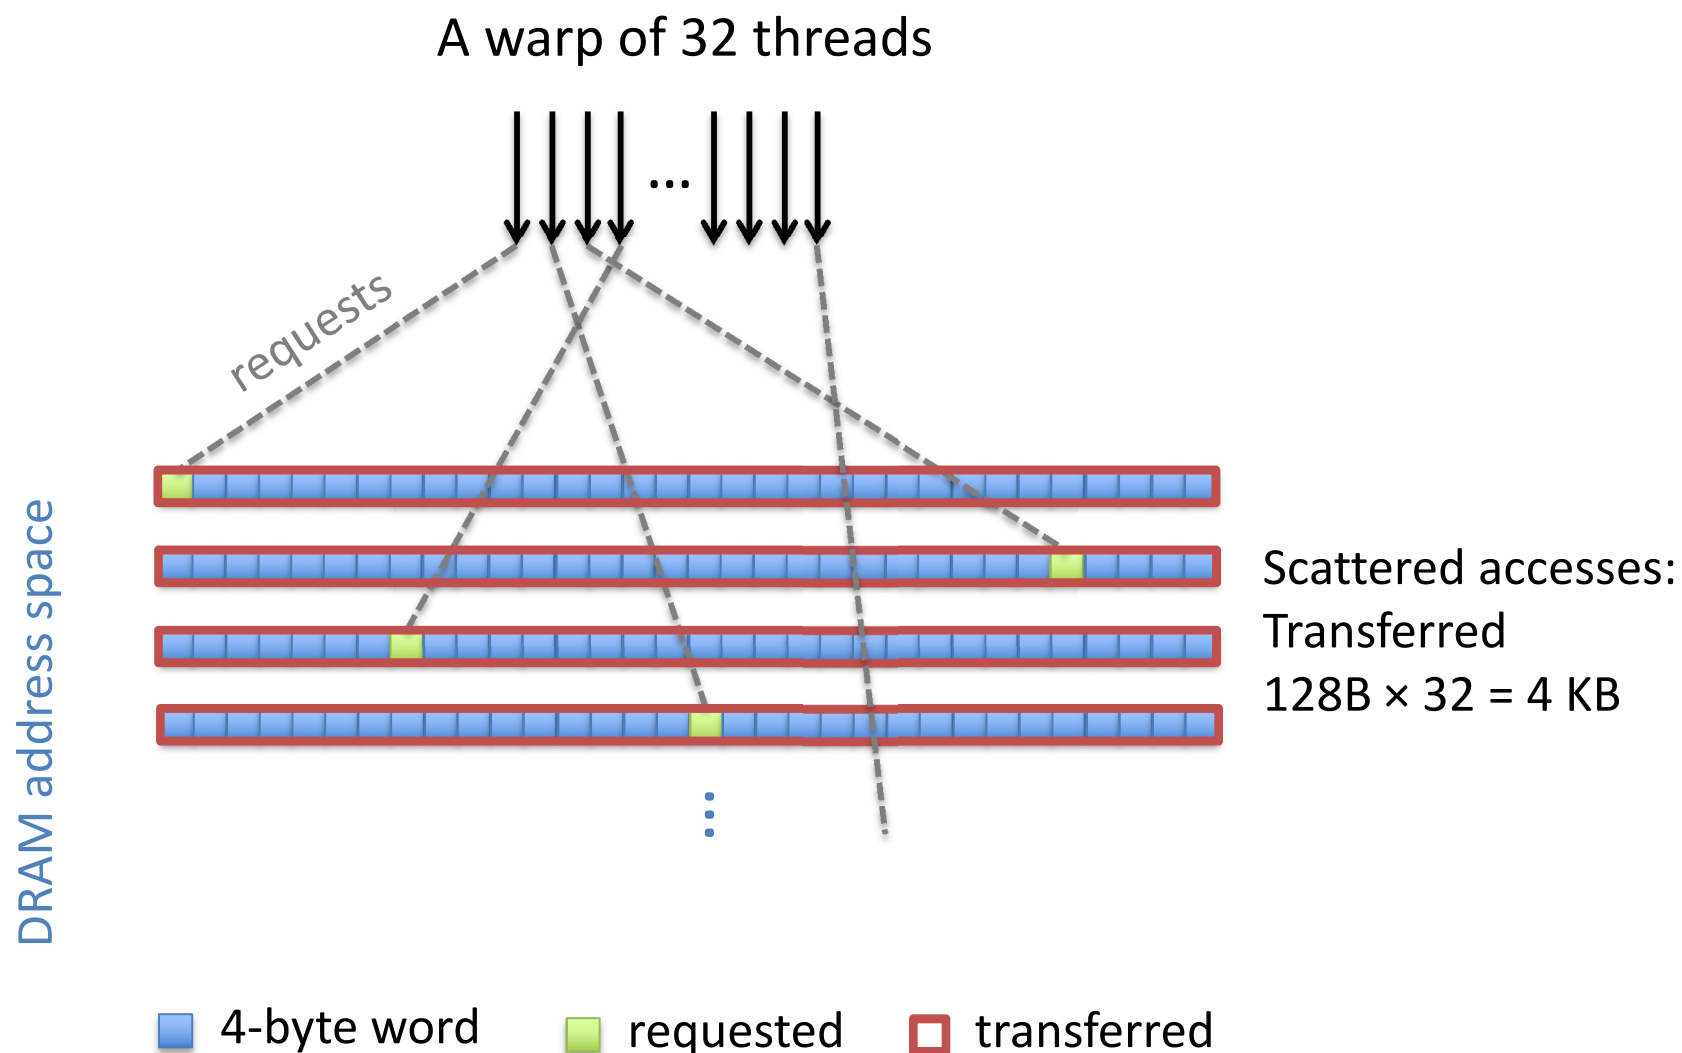


# Background: Effective Coalescing

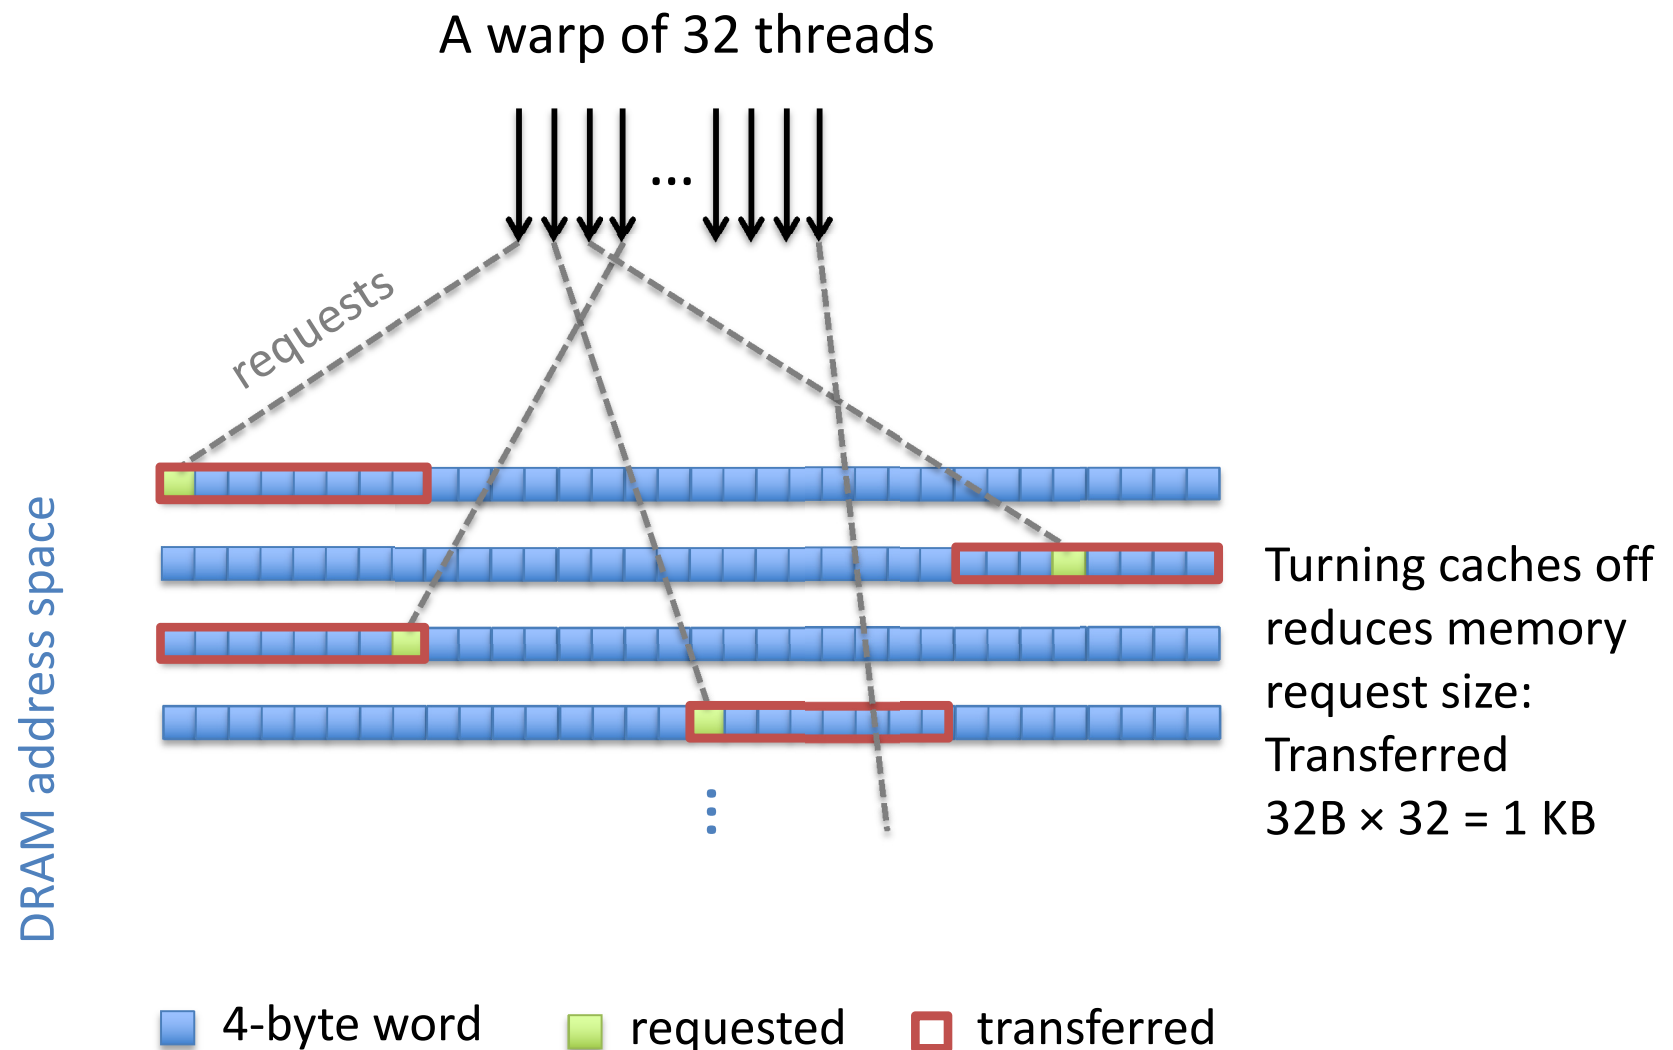






# Background: Poor Coalescing with Caching On



# Background: Poor Coalescing with Caching Off




# How to Predict Cache Utility?



Performance and cache hit rate are poorly correlated




# Global Memory Traffic Is a Better Indicator



L2-to-L1 read traffic better correlates with runtime changes




# Memory Access Locality Taxonomy




## Within-warp locality

Threads in the same warp access consecutive locations; coalesced, no caching needed



## Within-block locality

Threads in the same block access consecutive locations; short-term caching needed



## Cross-instruction reuse

The same thread accesses consecutive locations at different points in time; longer-term, difficult to cache



# Compiler Algorithm to Estimate Cache Utility

- Assumptions: Consider each instruction in isolation, control caching on a per-instruction basis, assume arrays have zero base addresses
- Step 1: Compute load addresses of load instructions
- Step 2: Estimate cache-on and cache-off traffic
- Step 3: Decide whether to cache for each instruction



# Example: Compute Load Addresses

- Use the constant folding/propagation framework to represent each thread's global memory load address as a function of its thread ID (tid), constants, and uncomputable ( $\infty$ )

```
// thread block size: 256
void toy (float *A, int *B, float *C, float *D) {
    int index = threadIdx.x;
    float x = A[index / 32];
    int y = B[index * 32];
    float z = C[y];
    D[index] = x * z;
}
```

## Results

index = tid  
load (tid / 32)  
load (tid \* 32)  
load ( $\infty$ )



# Example: Estimate Memory Traffic

- Goal: For an entire thread block, estimate the amount of global memory traffic when all threads execute one particular load instruction with and without caching

Load Inst	Expr	Requested Memory Addresses (Byte)				Cache-on Traffic	Cache-off Traffic	Cache?
		Warp 0	Warp 1	...	Warp 8			
A[index / 32]	tid / 32	0-3	4-7	...	28-31	128 B	$32 \times 8 = 256 \text{ B}$	Y
B[index * 32]	tid * 32	0-3 32-35 ...	1024-1027 ...	...	8196-8199 ...	$128 \times 256 = 32 \text{ KB}$	$32 \times 256 = 8 \text{ KB}$	N
C[y]	$\infty$							N

← 1 thread block →






# Caching Decisions

- For each load instruction, when an entire thread block executes it:
  - Cache-on traffic  $<$  cache-off traffic: cache for this load
  - Cache-on traffic  $>$  cache-off traffic: don't cache for this load
  - Cache-on traffic = cache-off traffic
    - Conservative strategy: turn caching off for this load
    - Aggressive strategy: turn caching on for this load
  - Unknown/uncomputable: likely to be scattered accesses based on our empirical observations, turn caching off for this load



# Cache On/Off Algorithm: Results



In most cases, our compiler algorithm matches or improves on best all-or-nothing choice



# Conclusions

- We should base GPU cache utility analysis on reducing global memory bandwidth requirement instead of hiding latency
- A memory access locality taxonomy helps analyze global memory traffic and consequently cache utility
- Automatable, static algorithm improves the average caching benefit from 5.8% to 18%



**THANK YOU!**

