A System-Level ISA and its Applications to Energy-Performance-Reliability Scheduling and Scratchpad Allocation

Yavuz Yetim, Wenhao Jia, Margaret Martonosi, Sharad Malik, Kelly Shaw*
Princeton University * University of Richmond

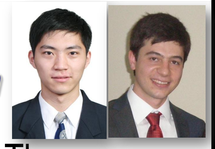

Platforms and Alternative Themes, Task #5.2.1 and 5.6.3

GSRC Annual Symposium
September 28, 2010 through October 1, 2010

Motivation

Problems:
- Rapidly increasing core counts
- Various heterogeneity (GPU, faster/slower cores)
- Decreasing reliability

Goal: Achieve performance, energy and reliability demands in future heterogeneous and dynamically changing multicore systems

SISA: Approach

Represent programs as graphs with app characteristics
- Data communication
- Length of computational tasks
- Reliability requirements
- Task dependency

Use SISA representation to do
- Static scheduling with Integer Linear Programming
- Pre-run resource mapping (such as scratchpad)
- Dynamic task management

SISA Graph

- No side effects inside a chunk
- Data flow marked explicitly

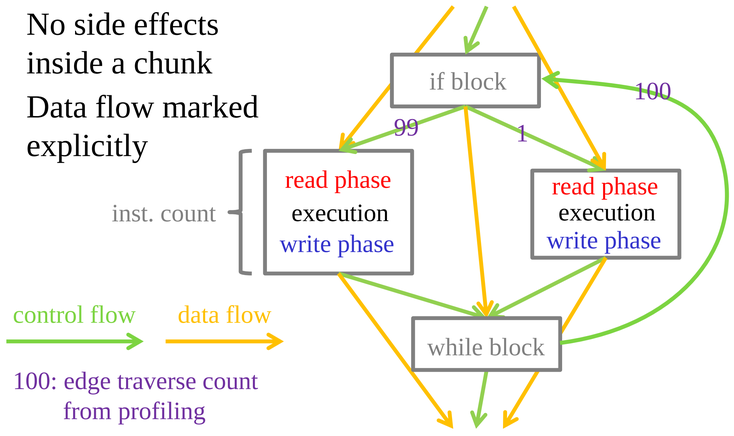

Application 1: Static Scheduling for Energy/Performance/Reliability

Novel Integer Linear Programming Formulation that
- ✓ optimizes periodic applications using pipelining
- ✓ includes communication overheads
- ✓ handles "lock" variables through mutual exclusion

Objective: Busy Energy + Idle Energy + Data Migration Overhead

Constraints:
- Overlap and Sequencing
- Error Tolerance
- Mutual Exclusion
- Pipelining
- Period and Deadline

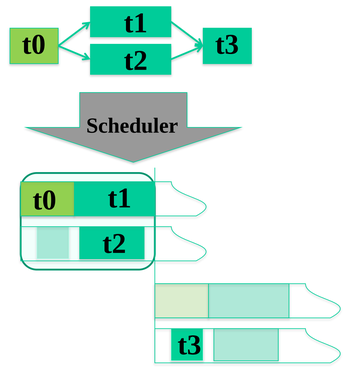

Results

Scheduler is run for StreamIt benchmarks on different configurations, which are assumed to be DVFS overscaled. Energy gains depend on:
- Hardware configuration
  - High idle power decreases gains
  - High error rate decreases gains for low error tolerance but increases gains for high error tolerance
- Application characteristics
  - Applications with workload-balanced tasks have better gains
  - Error tolerance of an application decreases energy significantly up to a point and flattens after all cores can be utilized

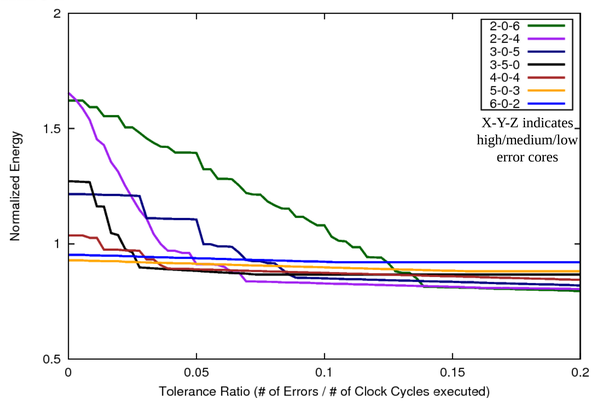

Application 2: Scratchpad Memory Allocation

Objective: Minimize program execution time on machines with local software-controlled memory

Method:
- Use the SISA graph representation of an application that has memory accesses and data flow marked
- Allocate the scratchpad for variables on the critical path through the Control Flow Graphs
- Allocate the scratchpad for the rest variables

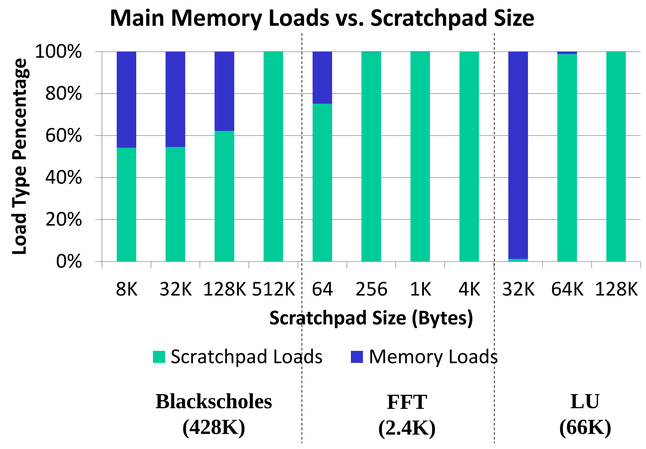

Results:

- On PARSEC benchmark suite
- Global memory loads reduce as available scratchpad space increases

Conclusions

- SISA graph-based program representation effectively exploits Performance/Energy/Reliability space
- Up to 34% energy savings can be achieved given application reliability requirements
- Based on data flow exposed by SISA graph, up to 99% memory loads can be eliminated with reasonable scratchpad size

Acknowledgements

The authors acknowledge the support of the Gigascale Systems Research Center, one of six research centers funded under the Focus Center Research Program, a Semiconductor Research Corporation program. In addition, this work was supported in part by National Science Foundation grant CCF-0916971.